

Parallel Implementation of LU decomposition

Adil Ansari

October 16, 2013

1 Basics

- Root directory contains three sub-directories namely 'Sequential', 'OpenMP' and 'MPI'.
- Each subdirectory has source code in the form of '*.c' file.
- Matrix is generated in a manner that it decomposes into a L and U containing only 1s and 0s.
- To submit jobs for various configurations run './submit.sh' on terminal. This will automatically submit all the jobs in the subdirectory to the **general-compute** queue of the ccr cluster.
- Outputs are generated in the /textbfoutput.txt file. *Sample outputs are included.*
- In the corresponding subdirectory run './plot.sh' on the linux terminal to generate a graphical visualization of the output. *gnuplot is required to generate graph.*
- Graphs are generated as '**Plot.pdf**'. *Please wait for the job run to finish and outputs to accumulate.*

2 Sequential Implementation

- Gaussian elimination algorithm was implemented that sequentially decomposes the square matrix.
- Algorithm was evaluated on input matrix size of 1000, 5000, 10000, 20000.
- Time taken to decompose the matrix grew exponentially with the increase in size.
- Since, this was a sequential implementation increase in compute nodes won't do anything.
- Since I was using Gaussian elimination that computes L and U matrices separately, I ran **out of memory** when matrix size of 50,000 was tried. This implementation makes two copies of the matrix of same size as input.




Figure 1: Sequential Decomposition Algorithm

3 OpenMP Implementation

- Gaussian elimination algorithm was implemented that uses the block wise decomposition in parallel.
- The **for** loops are parallelized in a manner that blocks of matrices are decomposed by dividing the work among parallel threads.
- Algorithm was evaluated for input matrix of sizes 1000, 5000, 10000, 20000 with a combination of 2, 4, 8, 16, 32 threads executing in parallel.
- On a fixed workload the decomposition was faster when more threads are executing in parallel. The execution was comparatively faster on larger workload due to the fact, parallelism was more effective.

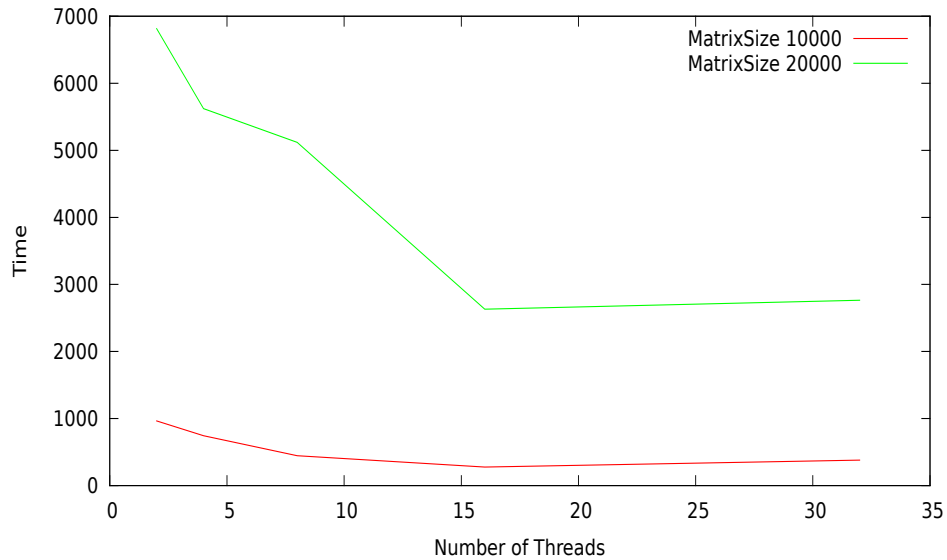


Figure 2: OpenMP Performance for fixed Matrix Size

- For a fixed number of cores the time increased exponentially with increase in matrix size.




Figure 3: OpenMP Performance for fixed Number of Cores

- The parallelism was ineffective on relatively smaller loads.
- Since I was using Gaussian elimination that computes L and U matrices separately, I ran **out of memory** when matrix size of 50,000 was tried. This implementation makes two copies of the matrix of same size as input.




Figure 4: OpenMP Decomposition Algorithm

4 MPI Implementation

- Cyclic distribution was used to accomplish LU factorization of the input square matrix.

- Each node is responsible for computing its own block and broadcast the result to rest of the nodes.
- Algorithm was evaluated for input matrix of sizes 1000, 5000, 10000 with a combination of 8, 16, 32 compute nodes working in parallel.
- For a fixed number of compute nodes the algorithm showed uniform behavior.




Figure 5: MPI Performance for fixed number of compute nodes

- For fixed workload the parallelism was more effective for larger workloads on maximum compute nodes.

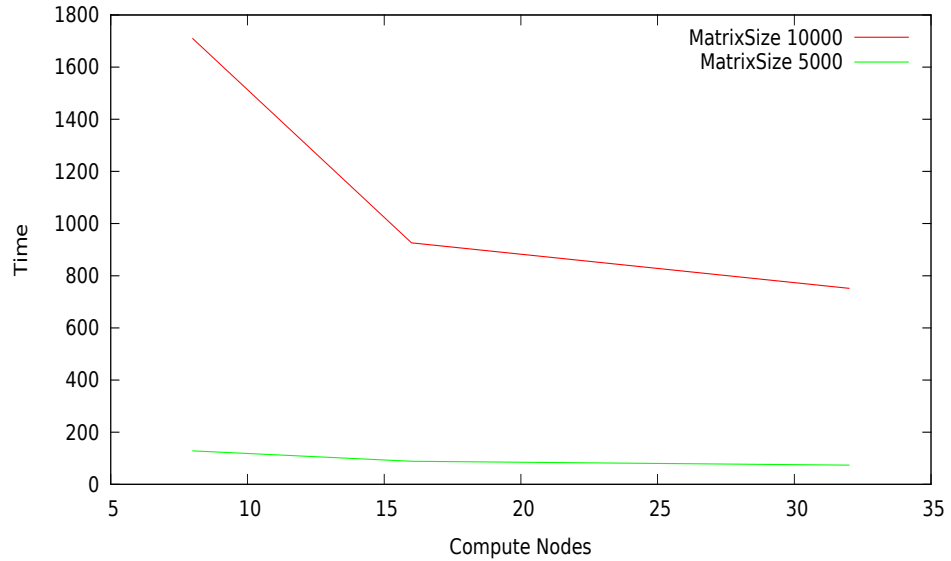


Figure 6: MPI Performance for fixed workload




Figure 7: MPI Decomposition Algorithm

5 Comparison

- Since, MPI involves communication overhead between different nodes, it was slower as compared to OpenMP.




Figure 8: OpenMP vs. MPI

- As expected sequential algorithm turns out to be the worst performer of the three.




Figure 9: Sequential vs. OpenMP vs. MPI

6 Scalability

LU factorization algorithm has a great extent of parallelization when scaled appropriately.